

Scalability & Future Plan

Date	29 September 2026
Team ID	SWTID-2026-3976
Project Name	AI FAQ Assistant API
Maximum Marks	3 Marks

Scalability & Future Plan

The team's original ER diagram and architecture diagram (see Project Documentation, EPIC 1 and ER Diagram Description) modeled Category, FAQ_TAG, TAG, FAQ_LIKE and ANSWER_FEEDBACK as fully separate, normalized collections, and a four-tier role system (Admin/Editor/Author/Reader) with draft-review-approve workflows, rather than the two-model, single-role implementation actually shipped this cycle. That original design doubles as the product roadmap:

Planned Change / Capability	Purpose	Status
Normalize Category into its own collection	Allow categories to be created, renamed and reported on independently of the FAQ enum they currently live in	Future roadmap (per original ER design)
Add TAG and FAQ_TAG (many-to-many tagging)	Let a single FAQ carry multiple free-form tags for finer-grained discovery than the fixed category enum	Future roadmap
Add FAQ_LIKE and ANSWER_FEEDBACK collections	Let users like an FAQ and mark an answer helpful/not helpful, feeding back into which FAQs surface first in search	Future roadmap
Introduce a role field and Admin-only routes	Enforce the Admin / Content Creator / Authenticated User / Public User tiers already described in Roles and Responsibilities, rather than treating every logged-in user identically	Future roadmap
Move from keyword search to true semantic search	Match customer questions by meaning rather than exact keyword overlap, closing the gap between the prose's "Semantic FAQ Search" feature and the current searchFAQs implementation	Future roadmap
Persist an AI-generation log	Record every Gemini-generated answer/FAQ pair (even ones the user doesn't save) for auditing and future prompt tuning	Future roadmap